

Introduction to Keras

Vincent Lepetit

Software Libraries for Deep Learning

All these libraries come with a Python interface:

Caffe: from Berkeley and now Facebook, plain text for model description. Mostly for computer vision. See also Caffe2

Theano: from University of Montréal. Discontinued in September 2017.

Torch/PyTorch: from Facebook.

TensorFlow: from Google.

...

Keras: from Google, built on top on TensorFlow or Theano, making it easy to build deep architectures.

Installing Keras

Install:

1. (on a Mac, MacPorts (or maybe HomeBrew));
2. Python 2.7 or Python 3.3+;
3. pip (utility software to install python libraries);
4. TensorFlow;
5. Keras;
6. matplotlib (Python graphic library), h5py (Python library required for loading models), Pillow (replaces PIL, Python library for loading and processing images);
7. jupyter

Please do these installations before the next lecture!

<https://www.labri.fr/~vlepetit/teaching>

jupyter notebook

The screenshot shows a web browser window with three tabs at the top: "regression", "mnist", and a blank tab. The URL bar shows "localhost:8888/notebooks/regression.ipynb#". Below the tabs is a toolbar with icons for Apps, Feature Visualization, Matrix Calculus, Emergence of grid-li..., deepmind/pycolab:..., and Other Bookmarks. The main content area is titled "jupyter regression" and shows three code cells:

```
In [1]: import numpy as np

def F(x1, x2):
    return np.sin(np.pi * x1 / 2.0) * np.cos(np.pi * x2 / 4.0)

In [2]: A = 2
nb_samples = 1000
X_train = np.random.uniform(-A, +A, (nb_samples, 2))
Y_train = np.vectorize(F)(X_train[:,0], X_train[:,1])

In [3]: from keras.models import Sequential
model = Sequential()

from keras.layers import Dense, Activation

nb_neurons = 20
model.add(Dense(nb_neurons, input_shape=(2,)))
```

Sending your Assignments

[IMPORTANT!!!]

1. Create an account on GitHub or equivalent if you don't have one;
2. Use jupyter to do the assignment. You should also add plain text to explain your solution. Put the jupyter file on your github account;
3. Send the link to your jupyter file on your github account by email to vincent.lepetit@u-bordeaux.fr with MVA-MP1, MVA-MP2, or MVA-MP3 and your last name and first name in the email subject.
Example:

Subject: MVA-MP1 Doe, John
Body: link to your jupyter file.ipynb

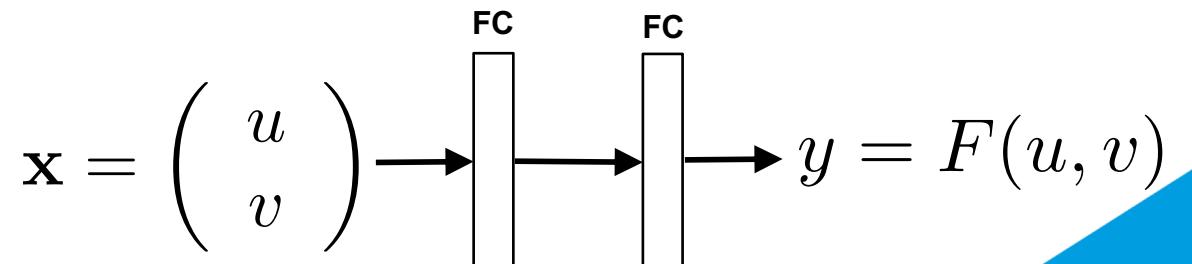
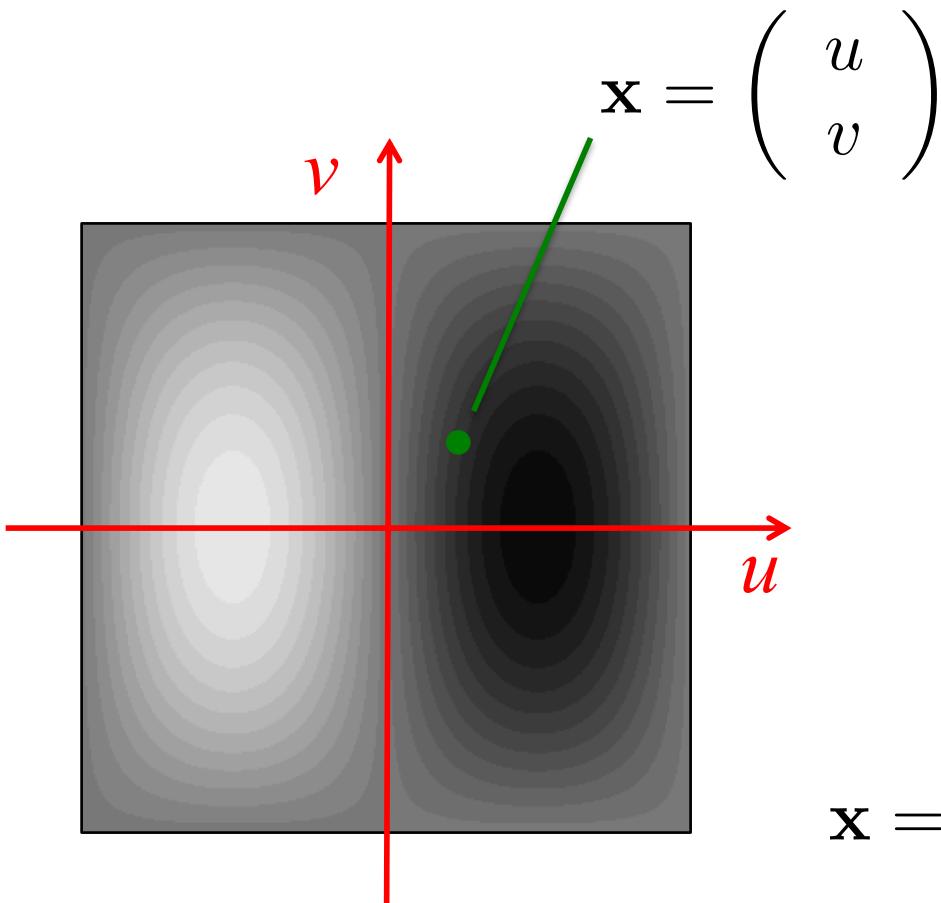
No exception!
No .zip, no .pdf!

4. No acknowledgment for receiving the email, sorry.

First Example: Two-Layer Network

A Two-Layer Network

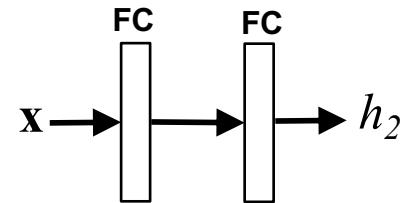
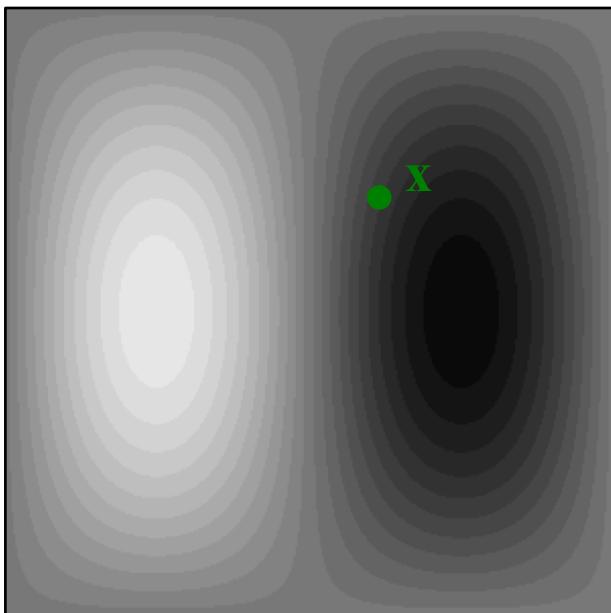
We will see how to train a two-layer network to approximate a 2D function $F(u, v)$:



Our Two-Layer Network

The input is a 2D point \mathbf{x} ;

The output is a scalar value h_2



Hidden layer:

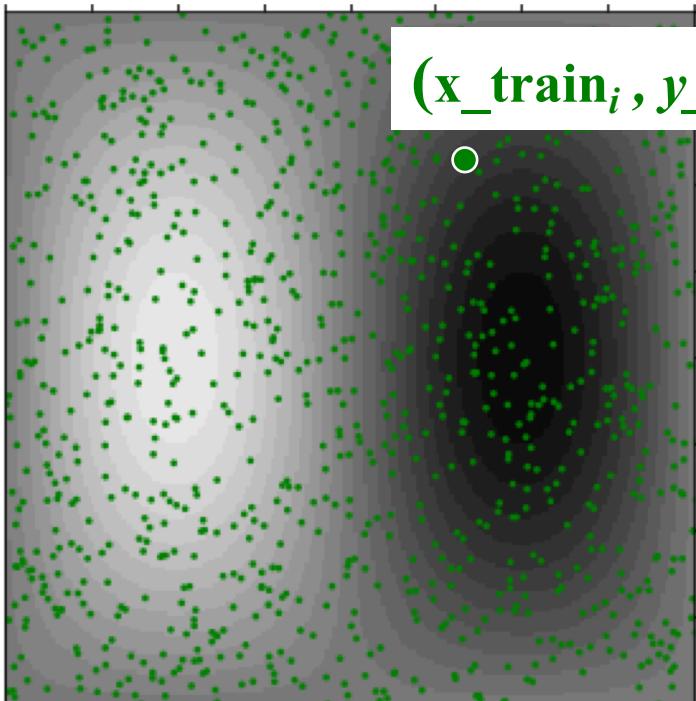
$$\mathbf{h}_1 = \text{ReLU}(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

Output layer:

$$h_2 = \mathbf{W}_2 \mathbf{h}_1 + b_2$$

Loss function

Training set:



Hidden layer:

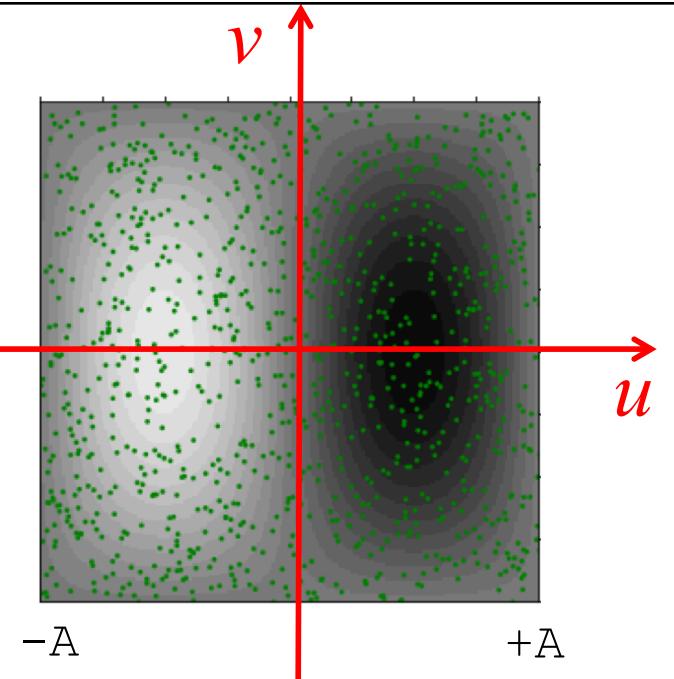
$$\mathbf{h}_1 = \text{ReLU}(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

Output layer:

$$h_2 = \mathbf{W}_2 \mathbf{h}_1 + b_2$$

$$\text{Loss} = \frac{1}{N_s} \sum_{i=1}^{N_s} (h_2(\mathbf{x}_{\text{train}}_i) - y_{\text{train}}_i)^2$$

Generating Training Data



```
import numpy as np

def F(x1, x2):
    return np.sin(np.pi * x1 / 2.0) * np.cos(np.pi * x2 / 4.0)

A = 2
nb_samples = 1000
X_train = np.random.uniform(-A, +A, (nb_samples, 2))
Y_train = np.vectorize(F)(X_train[:, 0], X_train[:, 1])
```

Models

In Keras, a deep architecture is called a *model*.

A model can be an arbitrary graph of layers.

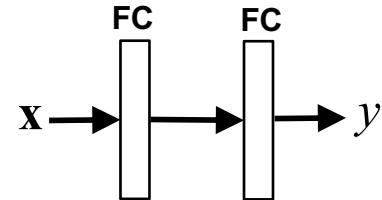
For this first example, we can use a Sequential model.

```
from keras.models import Sequential  
  
model = Sequential()
```

Defining the Network

Hidden layer: $\mathbf{h}_1 = \text{ReLU}(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$

Output layer: $h_2 = \mathbf{W}_2 \mathbf{h}_1 + b_2$



```
from keras.layers import Dense, Activation  
  
nb_neurons = 20  
model.add(Dense(nb_neurons, input_shape=(2,)))  
model.add(Activation('relu'))  
model.add(Dense(1))
```

Shortcut

```
from keras.models import Sequential  
from keras.layers import Dense, Activation  
  
nb_neurons = 20  
model = Sequential([  
    Dense(nb_neurons, input_shape=(2,)),  
    Activation('relu'),  
    Dense(1)])
```

Defining the Optimization Method

```
from keras.optimizers import SGD

sgd = SGD(lr=0.01,
           decay=1e-6, momentum=0.9,
           nesterov=True)

model.compile(loss='mean_squared_error',
               optimizer=sgd)
```

$$\text{Loss} = \frac{1}{N_s} \sum_{i=1}^{N_s} (h_2(\mathbf{x_train}_i) - y_train_i)^2$$

Running the Optimization

```
model.fit(X_train, Y_train, epochs=10, batch_size=32)
```

Output:

```
Epoch 1/10
1000/1000 [=====] - 0s 490us/step - loss: 0.0487
Epoch 2/10
1000/1000 [=====] - 0s 43us/step - loss: 0.0415
Epoch 3/10
1000/1000 [=====] - 0s 49us/step - loss: 0.0345
Epoch 4/10
1000/1000 [=====] - 0s 44us/step - loss: 0.0290
Epoch 5/10
1000/1000 [=====] - 0s 52us/step - loss: 0.0235
Epoch 6/10
1000/1000 [=====] - 0s 43us/step - loss: 0.0190
Epoch 7/10
1000/1000 [=====] - 0s 45us/step - loss: 0.0154
Epoch 8/10
1000/1000 [=====] - 0s 47us/step - loss: 0.0122
Epoch 9/10
1000/1000 [=====] - 0s 50us/step - loss: 0.0098
Epoch 10/10
1000/1000 [=====] - 0s 48us/step - loss: 0.0082
```

Prediction

```
x = [1.5, 0.5]
print(F(x[0], x[1]))  
  
x = np.array(x).reshape(1, 2)
print(x)
print(model.predict(x))
print(model.predict(x)[0][0])
```

Output:

```
0.6532814824381883
[[1.5, 0.5]]
[[0.5451795]]
0.5451795
```

Visualization

```
Width = 200
Height = 200
U = np.linspace(-A, +A, Width)
V = np.linspace(-A, +A, Height)
# Computes cartesian product between U and V:
UV = np.transpose([np.tile(U, len(V)), np.repeat(V, len(U))])
print(UV)
ys = model.predict(UV)
print(ys)
I = ys.reshape(Width, Height)
```

Output:

```
[[ -2.          -2.          ]
 [-1.9798995   -2.          ]
 ...
 [ 1.95979899  2.          ]
 [ 2.          2.          ]]

[[ 0.02076489]
 [-0.00082633]
 ...
 [-0.11296707]
 [-0.12041384]]
```

Visualization (2)

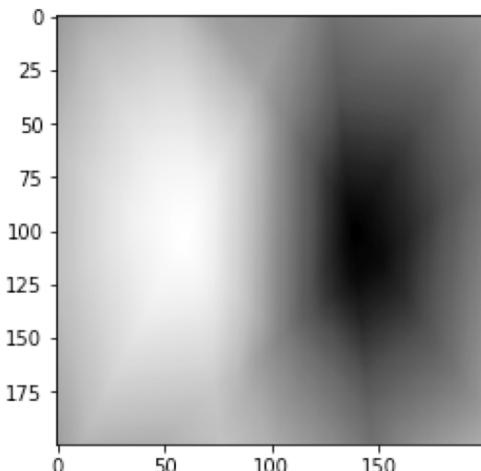
```
I = ys.reshape(Width, Height)

import matplotlib.pyplot as plt
import matplotlib.cm as cm

#Makes imageplotlib show the images inline
#in Jupyter notebooks:
%matplotlib inline
plt.imshow(I, cmap = cm.Greys)
```

Output:

```
<matplotlib.image.AxesImage at 0x117590390>
```



Second Example: CNN for MNIST

MNIST



Loading the Dataset

```
from keras.datasets import mnist  
  
(X_train, y_train), (X_test, y_test) = mnist.load_data()  
  
print( X_train.shape )
```

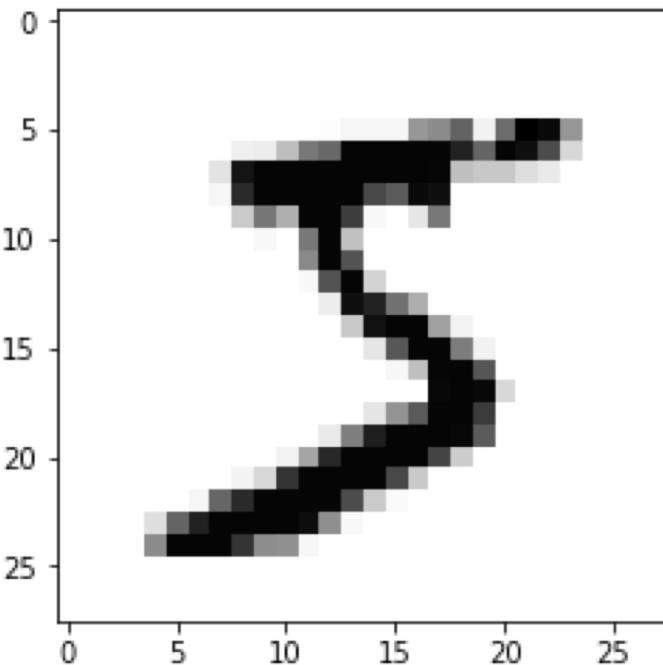
Output:

```
(60000, 28, 28)
```

Visualizing one Sample

```
import matplotlib.pyplot as plt  
import matplotlib.cm as cm  
%matplotlib inline  
plt.imshow(X_train[0], cmap = cm.Greys)
```

Output:



Reformatting the Input

```
X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)
X_train = X_train.astype('float32')
X_train /= 255
print(X_train.shape)

X_test = X_test.reshape(X_test.shape[0], 28, 28, 1)
X_test = X_test.astype('float32')
X_test /= 255
```

Output:

```
(60000, 28, 28, 1)
```

Reformatting the Desired Output

```
print(y_train.shape)
print(y_train[0:3])

from keras.utils import np_utils
y_train = np_utils.to_categorical(y_train, 10)
print(y_train[0])
```

Output:

```
(60000,)
[5 0 4]
[0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

Creating the Model

```
from keras.models import Sequential  
from keras.layers import Conv2D, MaxPooling2D  
  
model = Sequential()  
model.add(Conv2D(32, (3, 3), activation='relu',  
                input_shape=(28, 28, 1)))  
print(model.output_shape)
```

Output:

```
(None, 26, 26, 32)
```

Creating the Model (2)

```
model = Sequential()  
model.add(Conv2D(32, (3, 3), activation='relu',  
                input_shape=(28, 28, 1)))  
print(model.output_shape)  
  
model.add(MaxPooling2D(pool_size=(2, 2)))  
print(model.output_shape)  
  
from keras.layers import Dropout  
model.add(Dropout(0.25))  
print(model.output_shape)
```

Output:

```
(None, 26, 26, 32)  
(None, 13, 13, 32)  
(None, 13, 13, 32)
```

Creating the Model (3)

```
from keras.layers import Flatten  
model.add(Flatten())  
print(model.output_shape)  
  
from keras.layers import Dense  
model.add(Dense(128, activation='relu'))  
print(model.output_shape)  
  
model.add(Dropout(0.5))  
model.add(Dense(10, activation='softmax'))  
print(model.output_shape)
```

Output:

```
(None, 5408)  
(None, 128)  
(None, 10)
```

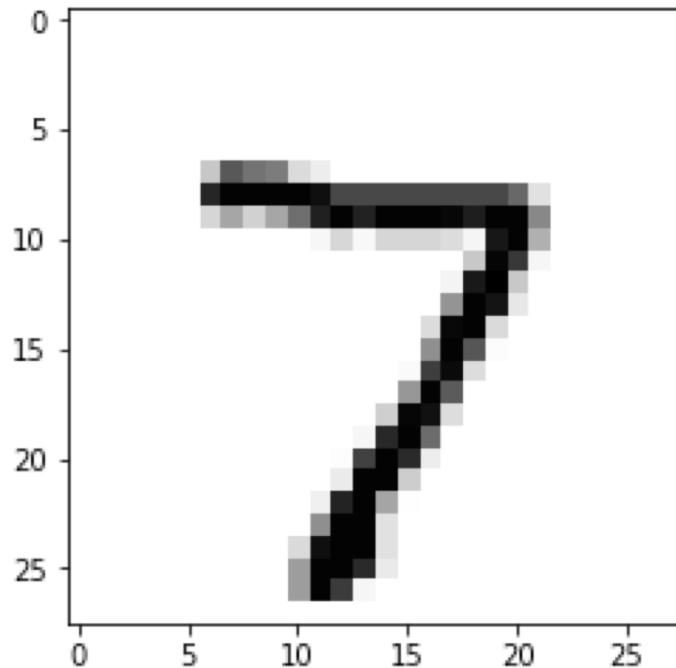
Optimization

```
model.compile(loss='categorical_crossentropy',  
              optimizer='adam',  
              metrics=['accuracy'])  
  
model.fit(X_train, Y_train,  
          batch_size=32, epochs=10, verbose=1)
```

Testing

```
plt.imshow(X_test[0].reshape(28,28), cmap = cm.Greys)  
  
print(model.predict(X_test[0].reshape(1, 28, 28, 1)))
```

Output:



```
[[2.7737193e-10 2.5943342e-08 2.5428611e-07 3.6613658e-06 1.0967714e-10  
 7.6563078e-10 1.0837641e-14 9.9999535e-01 2.9037880e-08 7.2273968e-07]]
```

Third Example:

Using VGG to Recognize Objects
in Images

Loading the VGG Model

```
from keras.applications.vgg16 import VGG16  
model = VGG16()  
print(model.summary())
```

Output:

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
etc.		

Loading the VGG Model

```
from keras.preprocessing.image import load_img  
# loads an image from file  
image = load_img('cat.jpg', target_size=(224, 224))  
image
```

Output:



Converting the Image

```
from keras.preprocessing.image import img_to_array  
# converts the image pixels to a numpy array:  
image = img_to_array(image)  
print(image.shape)  
image = image.reshape((1, image.shape[0], image.shape[1],  
                      image.shape[2]))  
print(image.shape)
```

Output:

```
(224, 224, 3)  
(1, 224, 224, 3)
```

Applying the Model

```
y_pred = model.predict(image)
print(y_pred.shape)

from keras.applications.vgg16 import decode_predictions
# convert the probabilities to class labels
labels_pred = decode_predictions(y_pred)
print(labels_pred)
```

Output:

```
(1, 1000)

[[('n02124075', 'Egyptian_cat', 0.5917598), ('n02123159',
'tiger_cat', 0.37640235), ('n02123045', 'tabby', 0.031016493),
('n02127052', 'lynx', 0.0006430329), ('n04589890',
>window_screen', 7.664625e-05)]]
```

Full Code

```
from keras.applications.vgg16 import VGG16
model = VGG16()
from keras.preprocessing.image import load_img
image = load_img('cat.jpg', target_size=(224, 224))
y_pred = model.predict(image)
from keras.applications.vgg16 import decode_predictions
labels_pred = decode_predictions(y_pred)
print(labels_pred)
```